

# Loan Default Detection



# The Problem

Financial institutions need to determine whether or not to approve someone's loan application

## Why is it worth solving?

Banks want to increase their profit margins while also reducing their exposure to risk



01

# Data Source

Credit Card Fraud Detection Dataset from Kaggle (<https://www.kaggle.com/mishra5001/credit-card>)



# Loan Default Detection Raw Dataset

1

## Shape

307,511 rows

2

## Variable Types

Integers, Strings, Numerical, Logical

3

## Predictors

	SK_ID_CURR	TARGET	NAME_CONTRACT_TYPE	CODE_GENDER	FLAG_OWN_CAR	FLAG_OWN_REALTY	CNT_CHILDREN	AMT_INCOME_TOTAL	AMT...
0	100002	1	Cash loans	M	N	Y	0	202500.0	
1	100003	0	Cash loans	F	N	N	0	270000.0	
2	100004	0	Revolving loans	M	Y	Y	0	67500.0	
3	100006	0	Cash loans	F	N	Y	0	135000.0	
4	100007	0	Cash loans	M	N	Y	0	121500.0	
...	...	...	...	...	...	...	...	...	...
307506	456251	0	Cash loans	M	N	N	0	157500.0	
307507	456252	0	Cash loans	F	N	Y	0	72000.0	
307508	456253	0	Cash loans	F	N	Y	0	153000.0	
307509	456254	1	Cash loans	F	N	Y	0	171000.0	
307510	456255	0	Cash loans	F	N	N	0	157500.0	


307511 rows x 122 columns




02

# Data Introduction

# Occupation and Loan Size



## Concentration of Loan Amount vs. Income



03

# Cleaning our Data



# Dropping Columns

- Drop repetitive **mode and median** columns that outline living standard
- Dropped outliers with income greater than or equal to **\$700,000**

Our resulting  
dataframe now has  
**306773 rows, 88  
columns**

# Adding Dummies

- Adding dummies for categorical columns using **pd.get\_dummies**

NAME_CONTRACT_TYPE_Cash loans	NAME_CONTRACT_TYPE_Revolving loans
1	0
1	0
0	1
1	0
1	0
1	0
1	0
1	0

Resulting dataframe

now has **306773 rows**,

**200 columns**



## Filling in NaN Values

- Tried **K-Means Imputer** but the process was very **slow and impractical** to our workflow because we had a large quantity of NaN values
- Instead, we calculated the **means** for each column and filled missing values based on the mean value

```
for c in df_with_dummies_drop_nan.columns:  
    df_with_dummies_drop_nan[c].fillna(value = df_with_dummies_drop_nan[c].mean(), inplace=True)  
df_with_dummies_drop_nan.isna().sum()
```

## Sampling Techniques/Split

- Randomly subsampled for 20,000 rows
- Experimented with **RandomUnderSampler & NearMiss** (undersampling) and **SMOTE** (oversampling)
  - **SMOTE** generates **synthetic** samples for **minority** class
  - **SMOTE** had relatively **poor** performance - samples are too artificial
- Next tried **NearMiss** and tested out different ratios for the **sampling\_strategy** parameter
  - We chose a **0.3** ratio between instances of default and no default
  - For every 3 defaults, there are 10 non-defaults
- After subsampling and undersampling, we have about **10,000 rows**
- **80/20** Train Test Split, Standardized and Normalized
- Based on how we sampled, the baseline has an accuracy of **70%**




04

# Descriptive Analysis & Predictive Models


# Logistic Regression

**Default LR**




**Training set accuracy: 80.90%**  
**Testing set accuracy: 82.77%**

**CV LR**



**Training set accuracy: 80.90%**  
**Testing set accuracy: 82.91%**

# Logistic Regression

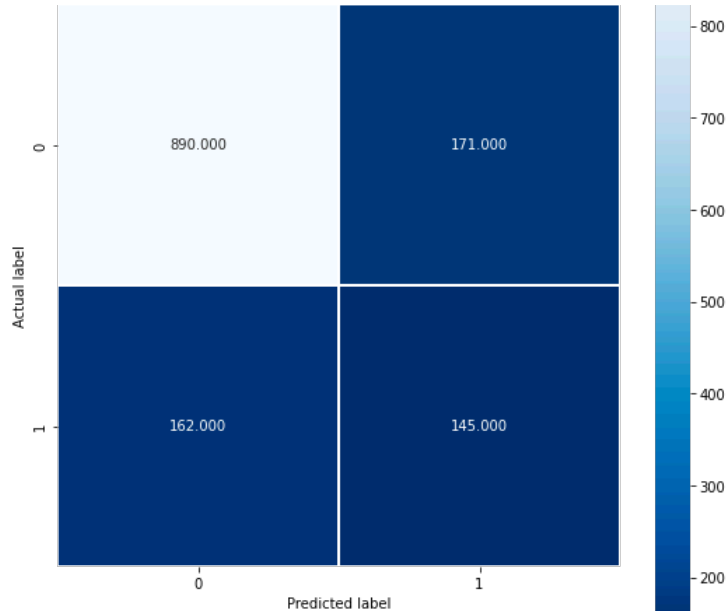


- The plot shows that the **higher the income**, the probability of **default increases**
  - We tried different parameters and all resulted in the same curve
  - We believe that this is due to **bias** within the dataset



# Bagging


## Default Bagging



Training set accuracy: 89.41%

Testing set accuracy: 73.69%

## CV Bagging

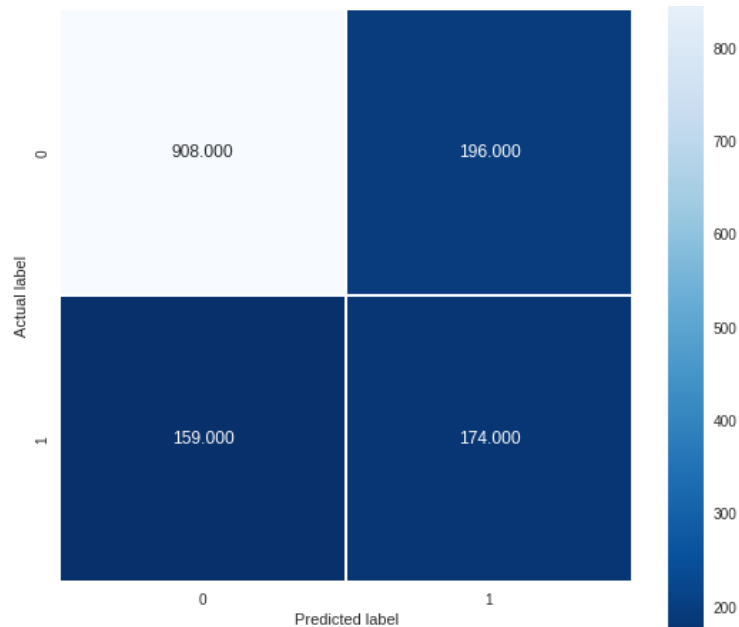


Training set accuracy: 88.72%

Testing set accuracy: 75.82%

# Boosting


## Default Boosting



Training set accuracy: 83.33%

Testing set accuracy: 79.36%

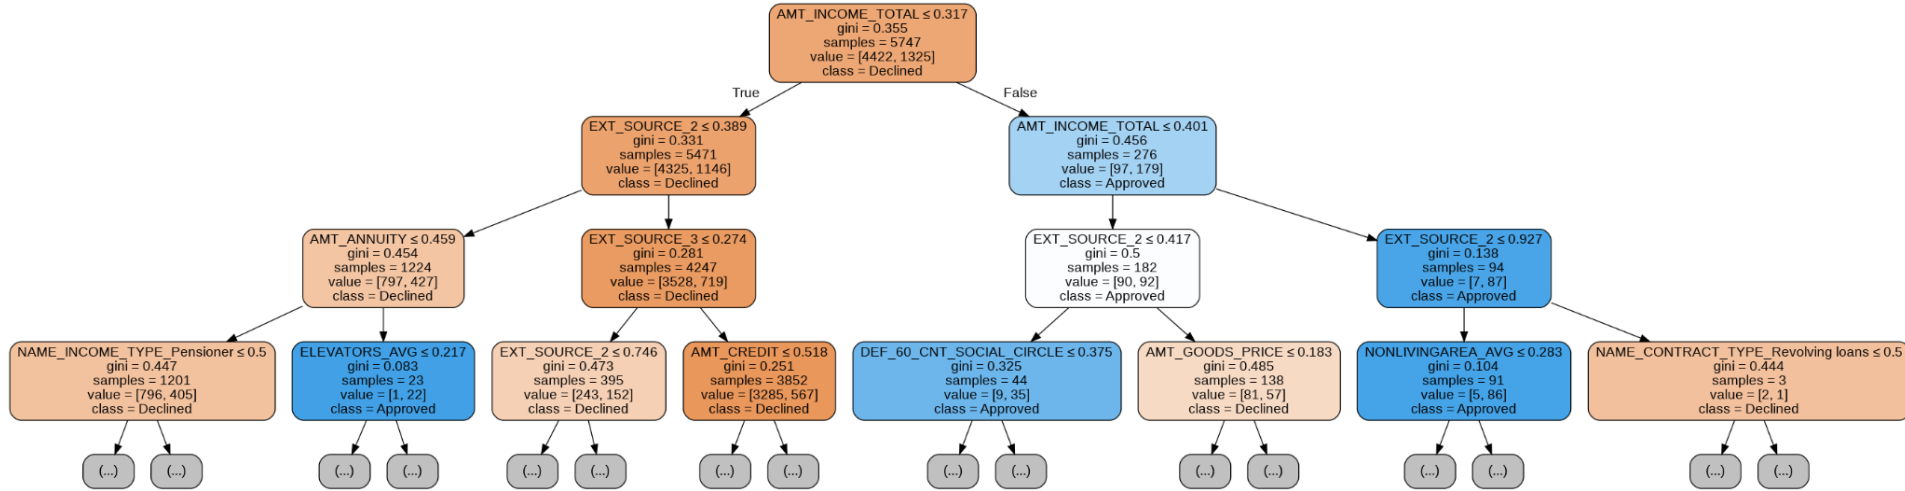
## CV Boosting




Training set accuracy: 83.74%

Testing set accuracy: 79.50%

# Decision Tree




# Decision Tree after CV (cropped)




# Decision Tree

## Default Decision Tree




Training set accuracy: 100%  
Testing set accuracy: 58.24%

## CV Decision Tree

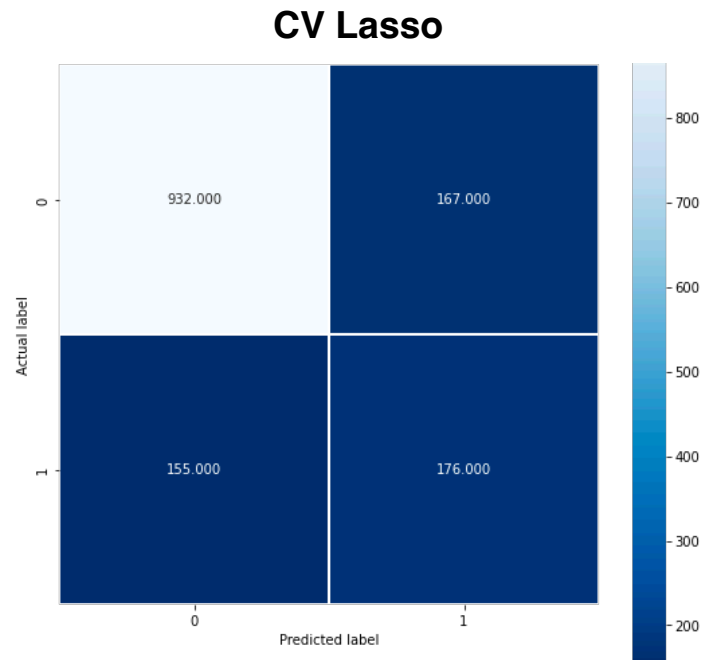


Training set accuracy: 79.26%  
Testing set accuracy: 77.52%

# Lasso




Training set accuracy: 76.48%  
Testing set accuracy: 78.65%



Training set accuracy: 80.79%  
Testing set accuracy: 82.48%

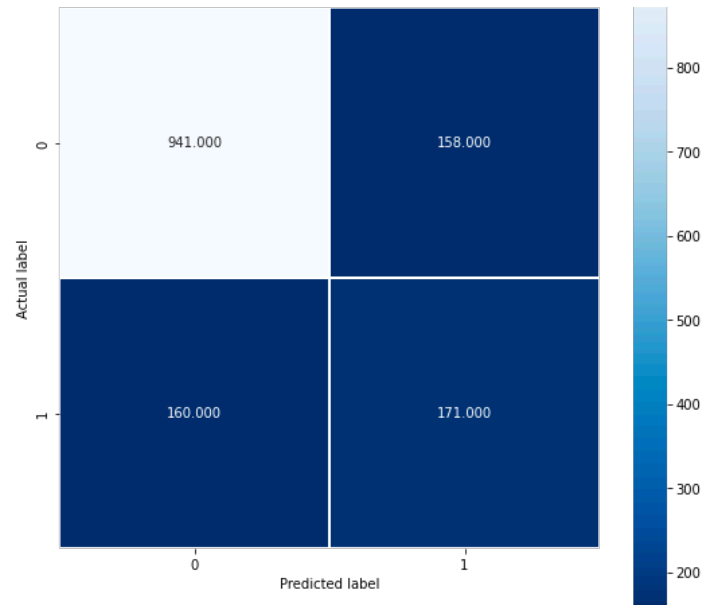
# Ridge

## Default Ridge



Training set accuracy: 80.95%  
Testing set accuracy: 82.26%


## CV Ridge



Training set accuracy: 80.99%  
Testing set accuracy: 82.12%


# Elastic Net

## Default Elastic Net



Training set accuracy: 79.98%  
Testing set accuracy: 82.20%

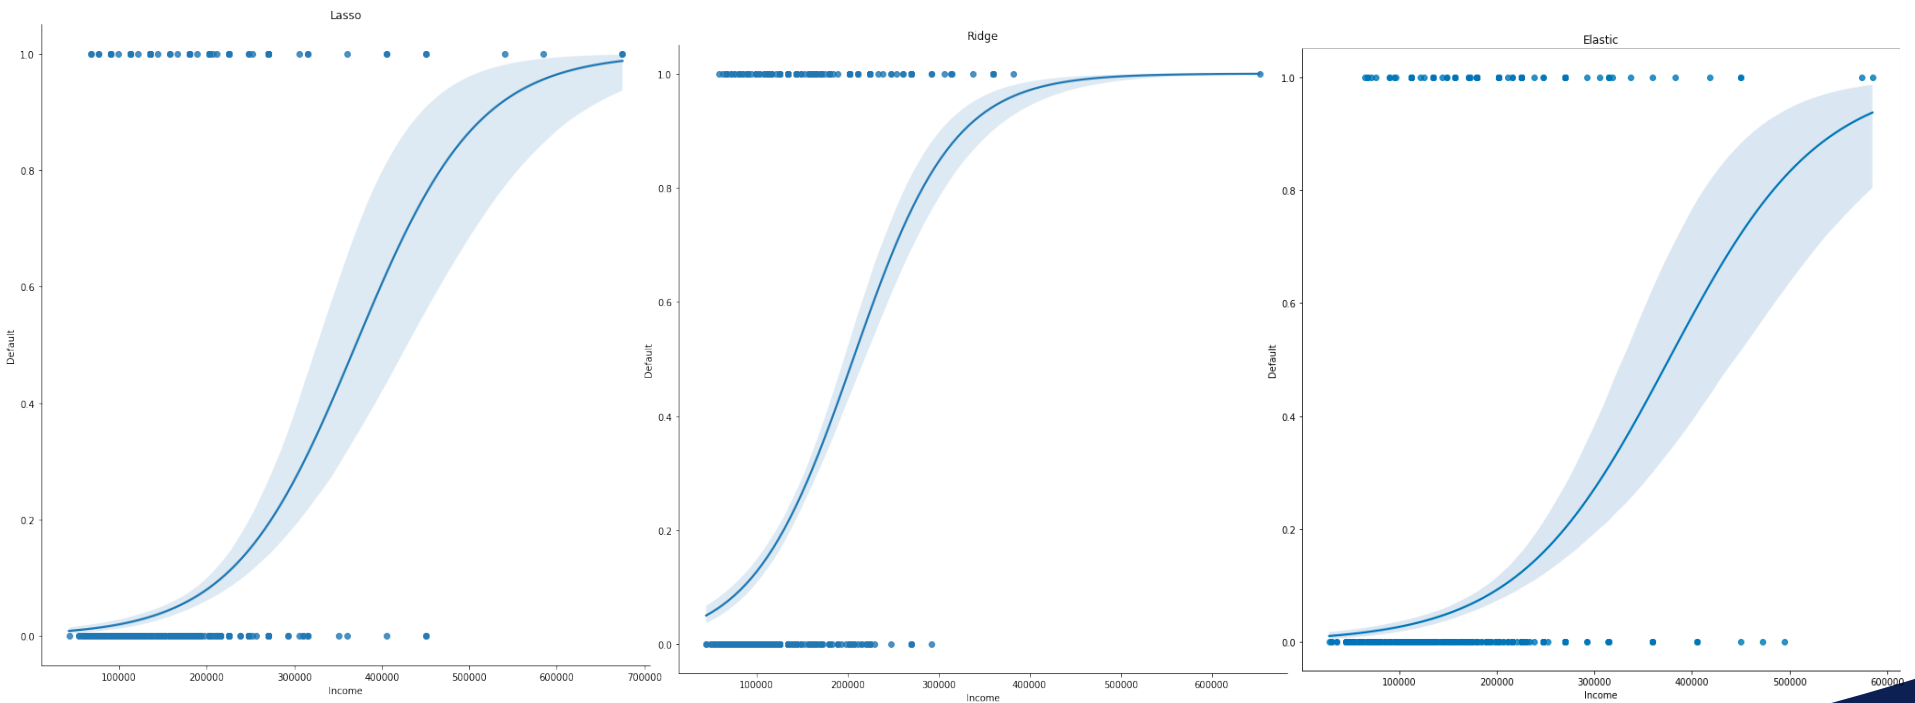
## CV Elastic Net



Training set accuracy: 79.86%  
Testing set accuracy: 81.50%




# Lasso, Ridge, Elastic Net



## Best Model based on accuracy

- **Lasso Net** is the most accurate model
- Why lasso has the highest accuracy
  - Lasso - some coefficients can become **zero** and **eliminate** the predictors from the model
  - Based on the heatmap, we have mostly columns that do not correlate to target



Training set accuracy: 80.79%

Testing set accuracy: 82.48%

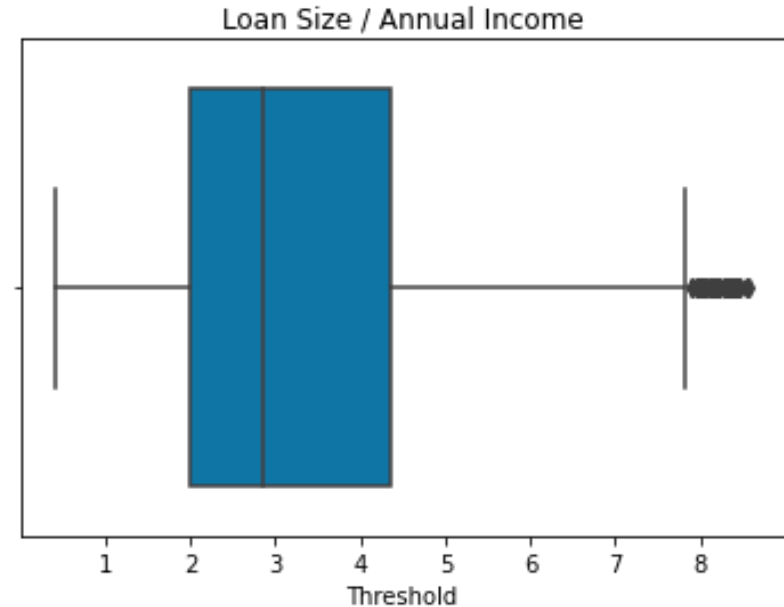
05

# Feature Engineering and Profit/Loss




## Risk Rating and Profit/Loss

- Interest rate is from **U.S. Treasury Yield**
- Inflation Rate of **2%**
- Removed outliers in terms of **loan size** and **income**
- Created **Risk Rating column** for each individual
  - Based on ratio of **Loan Size** to **Income**



# Feature Engineering Workflow


- Pass in each model's **predicted label** to our function
    - Using the predicted label of CV model
  - Calculate **profit** on instances where the predicted is **no default** and true label indicate **no default**
  - Calculate **loss** on instances where the prediction indicates **no default** but true label indicates **default**
  - Calculate **opportunity cost** on instances where the prediction indicates **default** but true label indicates **no default**
  - Aggregate **total profit and loss** to evaluate best model - the goal is to maximize profit
- 

## Profit/Loss Results

Model	Profit
Logistic Regression	\$179,461,405
Bagging	\$157,541,475
Boosting	\$155,338,172
<b>Lasso Regression</b>	<b>\$225,585,445</b>
Ridge Regression	\$189,043,931
Elastic Net	\$188,490,977
Decision Tree	\$212,958,162

# Lasso - Profit / Loss

CV Lasso  
Accuracy: 82.48%



- **Profit** - 932 people approved and did not default
- **Loss** - 155 people approved and defaulted
- **Opportunity cost** - 167 people rejected but would not default

06


# Obstacles



# Challenges

- Extremely large amount of features (**122 columns** before dummies) and dataset (**300,000+ rows**)
  - With a large dataset, running the models and functions was time consuming - **random subsampling** results in a new dataset every run
- In our dataset, the number of defaulting instances were very small compared to instances that did not default - **undersampling** using **NearMiss**
- Data was **inherently bias** given that it is collected from **multiple banks**
- Could not find **correlation** between most of the variables with the **target** variable

# Correlation Heatmap



Thank you